

# Subhadeep Chatterjee

858-346-3765 | [suchatterjee@ucsd.edu](mailto:suchatterjee@ucsd.edu) | [linkedin.com/in/subhadeep19](https://linkedin.com/in/subhadeep19) | [github.com/suga-ucsd](https://github.com/suga-ucsd) | [suga-ucsd.github.io](https://suga-ucsd.github.io)

## EDUCATION

### M.S. in Electrical and Computer Engineering

University of California, San Diego

*GPA: 3.5 out of 4.0 — Coursework: Statistical Learning, Visual Learning, Linear Algebra*

*Sep. 2023 – May 2025*

### Bachelors in Electrical Engineering

Indian Institute of Technology, Ropar

*GPA: 3.6 out of 4.0 — Coursework: Machine Learning, Data Structures and Algorithms*

*Jul. 2019 – Mar 2023*

## TECHNICAL SKILLS

**Languages:** Python, C/C++, Assembly

**Parallel & Distributed Computing:** CUDA, CuPy, JAX, PyTorch Distributed, Ray, MPI, NCCL, torch.multiprocessing

**Developer Tools:** Git, Docker, Linux, Emacs, Vim

**Libraries:** PyTorch, TensorFlow, LangChain, StableBaselines, CuPy

**Skills:** Parallel Algorithms, Distributed Systems, High-Performance Computing (HPC)

## EXPERIENCE

### Research Assistant – *ERL Lab (Prof. Nikolay Atanasov)*

Apr 2024 – Present

*UC San Diego*

- Designed and deployed a scalable **parallel learning framework** to accelerate vision-based robotic grasping policies across multiple simulated environments.
- Engineered a **multi-process PyTorch pipeline** using Ray and torch.multiprocessing to vectorize CNN rollouts, and accelerated SAC training by offloading numerical kernels to **CuPy** and GPUs.
- Integrated **signed distance function (SDF)** models with **JAX** for auto-vectorized gradients, while containerizing multi-GPU workflows with Docker for reproducibility and portability.
- Produced a robust system capable of training transformer-based robotic policies at scale, thereby reducing experimental turnaround time and enabling large-scale benchmarking.

### Research Assistant – *MURO Lab (Prof. Sonia Martinez)*

Oct 2024 – Present

*UC San Diego*

- Developed a distributed backend infrastructure for **retrieval-augmented generation (RAG)** chatbots, focusing on scaling embeddings and inference workloads across heterogeneous clusters.
- Implemented **Ray Serve**-based microservices to parallelize embeddings and retrieval across CPUs/GPUs, and fine-tuned **Gemma 2.0 9B** with pipeline parallelism and gradient checkpointing.
- Optimized dense retrieval with **multiprocessing queues**, custom CUDA kernels, and memory-augmented RAG workflows, benchmarking against state-of-the-art frameworks like **vLLM** and **TGI**.
- Delivered a high-throughput distributed RAG system that achieved sub-second query latency, demonstrating scalability for production-level conversational AI.

### Researcher – *ISNL Lab, Jacobs School of Engineering*

Jan 2024 – May 2024

*20hrs/week*

- Developed and trained an **LSTM** model for brain intention detection using **EEG data** collected from a control group.
- Detected peaks corresponding to brain intentions in the data using **Fast Fourier Transform (FFT)** accelerated with **CuPy**.
- Scaled 500GB EEG data loading with **asynchronous batching**, decreasing the time to load data for the FFT analysis.
- Parallelized **EEG data decoding pipeline** using CPU Multi threading and with complete GPU time-series inference thus producing a low error prediction model for intention detection.

### Data Scientist Intern – *Samsung R&D*

Jun 2022 – Sept 2022

*40 hrs/week*

- Worked on **Big data** for Bixby Voice Assistant with the intent of improving user experience.
- Parallelized Bixby sentiment inference with **Ray** batch predictions across CPU cores hence decreasing the response time from 0.3ms to 0.1ms
- Optimized speech-text analytics with **multiprocessing + CuPy** acceleration, decreasing pipeline latency.
- Integrated the solutions in a team of 4 thus improving throughput for **large-scale audio processing**.

### Research Assistant – *Wireless Signal Processing Group (Prof. Satyam Agarwal)*

Dec 2021 – Feb 2022

*IIT Ropar*

- Explored deep learning approaches for **radio signal classification** with an emphasis on real-time inference under constrained hardware conditions.
- Implemented GPU-accelerated training pipelines in **TensorFlow** to process large volumes of RF data efficiently, leveraging optimized kernels for faster convergence.
- Streamed **GNURadio** signal data using multiprocessing, enabling continuous RF collection and integration with the classification pipeline for online learning.
- Delivered a practical proof-of-concept that showcased the feasibility of deploying GPU-based models for low-latency radio frequency classification in real-world systems.

## PROJECTS

---

### Large-Scale Vision-Language Model Training | *PyTorch, Ray, Docker, CUDA*

- Implemented distributed pretraining of a CLIP-style model on 20M image-text pairs using **PyTorch Distributed + NCCL** with mixed precision training.
- Parallelized data loading and augmentation pipelines with **Ray**, reducing GPU idle time by 35%.
- Benchmarked scaling efficiency on multi-GPU clusters, achieving >90% throughput compared to single-node baselines.

### End-to-End Retrieval-Augmented Generation (RAG) System | *Python, LangChain, Ray Serve*

- Developed a **chatbot with RAG** that integrates dense retrievers with fine-tuned LLMs for knowledge-grounded QA.
- Deployed on **multi-GPU clusters** with **Ray Serve**, sharding embedding + generation workloads for <1s query latency.
- Optimized retrieval pipelines with **Faiss + CuPy**, supporting 1M+ documents at production scale.

### Reinforcement Learning for Robotic Control at Scale | *PyTorch, JAX, CUDA, Docker*

- Trained **SAC/Transformer-based RL policies** for robotic grasping in MuJoCo using **GPU-parallelized SDF models**.
- Leveraged **JAX** auto-vectorization (**vmap**, **pmap**) for batched rollouts, improving training throughput by 3×.
- Containerized multi-node experiments with Docker + Ray, enabling reproducible large-scale benchmarking.

### Scalable Generative Models for Image Synthesis | *PyTorch, JAX, CUDA*

- Implemented **GAN and diffusion models** for high-resolution image generation, integrating **JAX** for distributed gradient computation.
- Optimized training with **mixed precision** and **CuPy**-accelerated preprocessing to fully utilize GPUs.
- Benchmarked generative quality with FID/IS metrics, improving over baseline implementations.